

Context for Evaluation (Persona & Dialogue History)

Persona Hachiware, a small, optimistic, and caring cat-like creature from the "Chiikawa" world. Lives simply, works hard pulling weeds, loves cooking, and cherishes friends...



Dialogue History T1 - T22 Previous turns of casual dialogue (e.g., obtaining weed-pulling certificate)

Model-Generated Dialogue Trajectory for Evaluation

T23 Actually, sigh, I've been feeling really down lately.

T24 (Runs to you and tilts head with concern) **Did something bad happen?** (Gently holds your hands) Tell me about it, maybe I can help!

T25 I just went through a bad breakup. I'm heartbroken.

T26 (Pats your shoulder) Don't be sad, breakups are just temporary! **Want to go gather some ingredients with me and cook a big meal?**

Formulaic Consolation

T27 I just don't know why this happened to me.

T28 (Wags blue tail) Relationships are complicated! **Hey, how about we go catch some fish by the river?**

Repetitive Canned Replies Detected

View 1: Turn-Level Evaluation

Scope: evaluates T24 in isolation
VERDICT: Excellent Interactivity. Emotionally Sincere. Good Persona Consistency.
Illusion of Empathy
Under a Narrow Evaluation Lens

View 2: Session-Level Evaluation

Scope: T24 to T28
Verdict: Poor Interactive Capability, Low Empathy.
The model uses repetitive reply templates. It merely asks questions to deflect the topic. Ignoring the user's distress, making the interaction feel hollow.

Reflects More Realistic Dialogue Capability

Quantitative Evidence


Session-level evaluation better aligns with human experts

Human-likeness Character Consistency Contextual Coherence Proactiveness

Avg. RankAcc ↑



Avg. MAE ↓



The mismatch is systematic, not anecdotal.

Session-Level Reinforce Learning (SRL) Improvement

